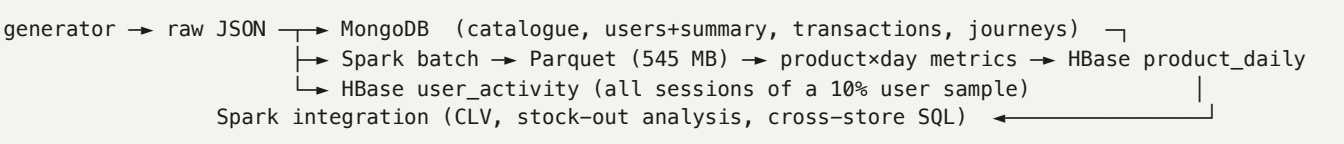
Distributed Multi-Model Analytics for E-commerce Data

MSDA9215 Big Data Analytics: Final Project · Patrick Niyonshuti (ID 101420) · November 2026
Code and results: github.com/Patrick101420/auca-msda9215-ecommerce-analytics

1. Overview and system architecture

The project builds an analytics system for a synthetic e-commerce platform over 90 days: 10,000 users, 5,000 products in 25 categories, 2,000,000 browsing sessions (20.0M page views) and 500,000 transactions (4.0 GB of JSON). Each technology is given the job its data model is good at:

- **MongoDB (document)** is the operational store: the product catalogue, user profiles with an embedded purchase summary, transactions with embedded line items, and complete session journeys. It serves point lookups and entity-centred aggregations.
- **HBase (wide-column)** is the time-series serving layer: a per-user activity stream keyed for newest-first retrieval, and a per-product daily performance series keyed for time-range scans.
- **Apache Spark** (PySpark 4.2, local mode, 12 cores) is the analytical engine. It cleans the raw JSON into Parquet, computes the product-day metrics that are bulk-loaded into HBase, runs the recommendation and cohort jobs and Spark SQL, and joins data pulled from MongoDB and HBase for the integrated analyses.



Infrastructure ran in Docker (MongoDB 7.0, HBase 2.1 with Thrift) on one laptop (15 cores, 24 GB). The course generator was adapted in one respect: it kept all 2M sessions in a Python list, which needs well over 20 GB of RAM. The adapted version writes each 100k-session chunk as it fills, which produced identical files with a measured peak memory of 1.04 GB (191 s).

2. Data quality: what the data can and cannot support

Before any analysis, job `01_clean_normalize.py` counted rule violations instead of silently fixing them (`output/dq_report.json`). Several of these findings shape the conclusions later.

Check	Result	Consequence
Session-ID collisions	2	IDs are `uuid4().hex[:10]` (40 bits); the birthday bound predicts ≈1.8 at 2M. The two pairs belong to different users, so the natural key is (session_id, user_id). MongoDB's unique `_id` rejected both; the HBase row key (which contains user_id) was unaffected.
“Converted” sessions with no transaction	149,276	Once the 500k transaction target is reached (≈1.25M sessions in), the generator still labels sessions converted. The *label* overstates conversion: 20.3% labelled vs 12.8% with a recorded order.
Orders not linked to any session	243,389 (48.7%)	These carry 57.6% of revenue ($268.7M of $466.6M) with a higher order value ($1,104 vs $771). Journey-based attribution can explain less than half of revenue.
Line items sold at a price ≠ list price valid at sale time	27.9%	The generator stamps today's price on past sales (net effect +$170,914). Price-elasticity analysis is not possible.
Products with `subcategory_id`	0 of 5,000	The brief's samples show it; the generator never writes it. Margins are therefore estimated at category level (mean of subcategory margins).
Users: non-US country with a US state	9,952	Geography is internally inconsistent, so it is used only as a label, never for spatial analysis.
Arithmetic: subtotal = Σ items, total = subtotal − discount	0 violations	Monetary fields are trustworthy.

3. Data modelling (Part 1)

3.1 MongoDB: model around the entity that is read together

Collection	Design	Justification
`categories` (25)	Subcategories **embedded**	Bounded (3–5) and always read with the parent.
`products` (5k)	Category id + name copied in (**extended reference**); `price_history` embedded	Catalogue pages need the category name without a `$lookup`; category names rarely change. Price history is bounded (1–3).
`users` (10k)	Profile + **computed** `purchase_summary` (orders, spend, AOV, first/last purchase) + **subset** of the 5 latest orders	Built by an aggregation with `$merge`. Profile pages and segmentation read one document. Embedding *all* orders would grow without bound (16 MB document limit).
`transactions` (500k)	Line items **embedded**, each with a product-name/category snapshot	An order is created and read atomically. The snapshot also preserves what the customer saw if the catalogue later changes.
`sessions` (2M)	Whole journey embedded; cart stored as an *array*, not a dict keyed by product id	One read = one journey. Field names that are data values cannot be indexed or unwound.

Indexes follow the queries: `{user_id, timestamp:-1}` and a multikey `items.product_id` on transactions; `{user_id, start_time:-1}` and `start_time` on sessions. `explain()` confirms IXSCAN with 10 keys/10 documents examined for "a user's recent orders" (1 ms). Load rates: transactions 56k docs/s, sessions ≈30k docs/s. WiredTiger compression stored 3.2 GB of session data in 891 MB.

Aggregation pipelines (`mongodb/aggregations.py`): **P1** monthly revenue by category with each category's top-3 products (`$unwind` → `$group` → `$sort` → `$group/$push` → `$slice`, 2.2 s, no `$lookup` thanks to the extended reference). **P2** spend-quintile segmentation from the embedded summary only (17 ms). A first attempt with fixed RFM cut-offs gave 47 / 2,527 / 7,424 / 2 users because every synthetic user orders ≈50 times in 90 days, so `$bucketAuto` was used to let the data set the boundaries. The top quintile spends more through *both* more orders (58 vs 42) and a higher AOV ($1,007 vs $855). **P3** most-viewed vs best-selling products over 30 days. Rewriting the `$lookup` from `$expr/$in` (cannot use the multikey index) to `localField/foreignField` cut it from 30.9 s to 19.7 s. Profiling showed the remaining 17.8 s is unwinding 664k journeys (6.6M page views). This is the document model's weak spot: a question *across* many journeys rather than *within* one.

3.2 HBase: design the row key for the scan

Table	Row key	Column families	Why
`user_activity`	`user_id#(9999999999−epoch)#session_id`	`s` summary (9 cells), `pv` page views (sparse, qualifier = sequence no.), `c` cart (sparse, qualifier = product_id)	Rows sort lexicographically, so a prefix scan returns a user's sessions **newest first** and `LIMIT n` stops early. Each family is its own HStore, so listing sessions reads only `s`, never the larger `pv` store.
`product_daily`	`product_id#yyyymmdd`	`m` counters, `dv` views by device, `rf` views by referrer (sparse)	A product's history is one contiguous key range: `STARTROW prod#20261101, STOPROW prod#20261201`. Zero counters are not stored at all.

Both tables are pre-split into 5 regions on key boundaries, with ROW bloom filters and `VERSIONS=1` (metrics are recomputed and overwritten by Spark). For `user_activity` we sampled **users, not sessions**: all 200,080 sessions (4.0M cells) of every 10th user. Sampling 10% of sessions would leave *every* user's stream 90% incomplete, which would make the per-user query wrong for everyone. `product_daily` holds all 421,769 product-days, loaded from Spark's output at 25.6k rows/s. Queries (`hbase/queries.py`, shell equivalents in `queries.hbase`):

Query	HBase mechanism	Rows	ms
Last 5 sessions of user_000420	prefix scan, family `s`, LIMIT 5	5	160 (cold) / 3.5 (warm)

Sessions in a 14-day window (range scan)	start/stop rows from reversed timestamps	26	39
One session's click-stream	`get`, families `pv`, `c`	1	41
Converted sessions (value scan)	SingleColumnValueFilter on `s:status` (server side)	35	176
Mobile sessions of users 100–199 (hybrid)	row range + family + value filter	674	152
One product, one month	range scan on `product#yyyymm`	11	31

3.3 Why each piece of data lives where it does

Entities that are created, updated and read as a whole (an order, a product, a profile, one journey) go to MongoDB, whose flexible schema and secondary indexes suit them. Append-only event streams, read by key range and time and sparse by nature, go to HBase: its sorted row keys make “latest N” and “between dates” a contiguous read, and absent cells cost nothing. Neither is used for whole-dataset analytics. That belongs to Spark over columnar Parquet. In CAP terms both stores are CP as deployed (single primary / single RegionServer per region). For a shop this is acceptable: a briefly unavailable recommendation panel is better than a wrong stock level.

4. Spark processing (Part 2)

4.1 Batch pipeline and optimisation choices

1. **Clean and normalise** (91 s for 4.0 GB): **explicit schemas** instead of inference. Inference costs a full extra pass, and because `cart_contents` is a dict keyed by product id, it would have produced a struct with thousands of fields. Declared as `MapType` it is one column that `explode()` unpacks. Nested arrays become normalised tables via `posexplode`. Timestamps are typed; empty strings become NULL; duplicates are removed on the composite key. The parsed sessions are `persist(MEMORY_AND_DISK)`-ed because they feed three writes plus the quality checks: a **branching lineage**, which is exactly when caching pays (Lecture 4). Output is Parquet (7.3× smaller than the JSON), with `page_views` **partitioned by event_date** for pruning.
2. **Product×day metrics**: conditional sums produce the sparse device/referrer breakdowns in a single pass (no `pivot()` discovery job). The result is bulk-loaded into HBase.
3. **Recommendations**: co-viewed pairs with the RDD API (`flatMap` → pairs, `reduceByKey`, which combines map-side before the shuffle, unlike `groupByKey`) and the per-product view counts **broadcast**. Each pair gets support, confidence and **lift**. “Bought together” pairs use a DataFrame self-join with broadcast dimension tables.
4. **Cohort analysis** by registration month, with a window `lag` to show change over time.

Engine settings: `local[12]`, 10 GB driver, 48 shuffle partitions (the default 200 would create tiny tasks at this scale) and Adaptive Query Execution. One pitfall worth recording: passing a naive Python `datetime` into a Spark filter silently converted it from the OS zone (UTC+3), shifting a 30-day window by 3 hours. Bounds are now passed as strings and cast inside Spark.

4.2 Spark SQL

`05_spark_sql.py` registers the Parquet tables as views and runs: a device × referrer funnel (CTE with conditional aggregation); the best product per category (`RANK() OVER (PARTITION BY category)` plus share of category); a **price-at-time-of-sale audit** (non-equi join to `price_history` with `ROW_NUMBER` to pick the price valid at the order timestamp); inter-purchase intervals with `LAG` (median 1.22 days); and weekly revenue with an estimated gross profit. For data that “lives” in the operational stores, MongoDB users (with the embedded summary, 10k docs in 0.15 s) and the HBase `product_daily` series (421k rows in 4.6 s) are loaded into DataFrames, registered as views, and joined in a single SQL statement (country value × each country's favourite product × that product's HBase funnel). Reads go through the native drivers with work pushed down to each store. The Mongo-Spark and HBase-Spark connectors would parallelise these reads across executors, which matters at production scale, but at these result sizes a driver-side read costs nothing measurable.

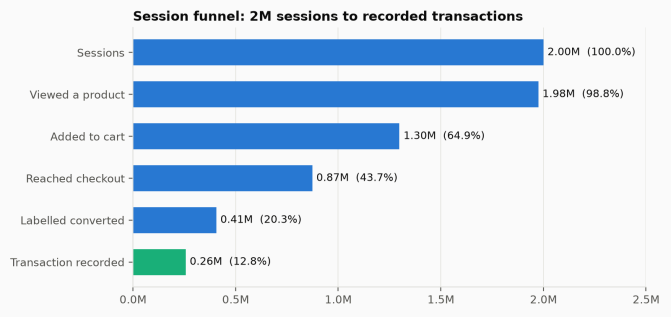


Fig. 1. Funnel. Conversion is 20.0–20.5% for every device × referrer pair (`output/sql_funnel_device_referrer.csv`): channel makes no difference.

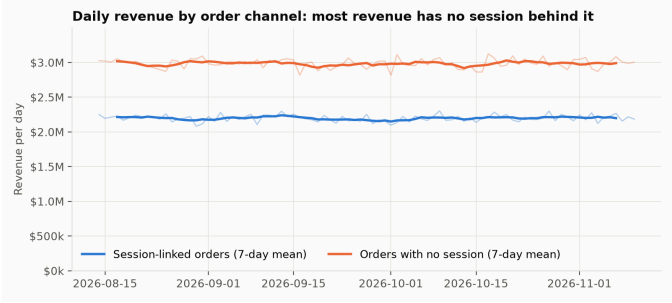


Fig. 2. Daily revenue. Orders with no session behind them earn more than session-linked orders on every day.

5. Integrated analytics (Part 3)

5.1 Customer lifetime value: does browsing engagement add anything?

Business question. Which customers will be worth most next quarter, and does on-site engagement predict future spend *beyond* purchase history? If it does not, a nightly join of 2M sessions to the customer table can be dropped. **Sources.** MongoDB transactions (one pushed-down aggregation splits each user's history at a cutoff, 0.46 s); MongoDB users; HBase `user_activity`, family `s` only, with a server-side `SingleColumnValueFilter` on `s:start < cutoff` (133,770 rows, 2.6 s; ISO timestamps compare correctly as bytes); Parquet line items → category-weighted margin per customer. **Method.** The 90 days are split into 60 days of calibration and a 30-day holdout. Spark MLlib linear regressions on the 1,000 HBase-sampled users (70/30 user split): M1 = history (frequency, monetary, recency, tenure), M2 = M1 + engagement (sessions, minutes, page views, converted sessions). CLV(12 m) = predicted 30-day spend × 12 × estimated margin.

Model	Test R^2	Test MAE	RMSE
Naive: next 30 days = ½ × last 60 days	–	$4,357	–
M1 history	−0.010	$3,523	$4,305
M2 history + HBase engagement	−0.018	$3,527	$4,323

Result. Neither model explains future spend ($R^2 \approx 0$; every feature's correlation with holdout spend is $|r| < 0.04$). The models beat the naive rule only by predicting the population mean. The top predicted decile delivered 10.0% of actual holdout spend, exactly what chance gives (Fig. 7b). This is a property of the data, not a modelling failure: the generator assigns every session and order to `random.choice(users)`, so past behaviour carries no information about the future. Without the holdout design, the same pipeline would have produced confident-looking CLV tiers that are pure noise.

5.2 Where is demand lost: stock-outs or weak conversion?

Business question. How much revenue was lost to products running out, and which in-stock products under-convert? **Sources.** HBase `product_daily` (views, cart adds, units per product per day); MongoDB `products` (stock, active flag, price, category); Spark for joins and per-category medians. The method took three attempts, and the rejected ones are part of the result:

1. **v1, units-per-view vs category median, all products: $125.9M "leaked". Rejected.** Half the catalogue (2,492 products) sold out, and the site almost stops showing a product once its stock is zero. Sold-out products keep the sales they made but lose views, so their units-per-view is inflated (0.46 vs 0.22). That raises every category median and makes in-stock products look leaky. The confounder is inventory.
2. **v2, an event study around each product's sell-out day. Rejected.** Views of sold-out products are spread evenly across the 90 days (second-half/first-half ratio 1.02, identical to in-stock products), and the computed sell-out days cluster at the end of the window. The generator depletes stock in *generation* order but stamps

each session with a *random* date, so the inventory history is inconsistent with the data's own clock. A time-series method cannot work here.

3. **v3, a cross-sectional comparison within category (adopted).** Missing views = median views of in-stock peers − own views. Lost units = missing views × the peers' median units-per-view. Lost revenue = lost units × price. Conversion leakage is measured only among in-stock products against in-stock medians.

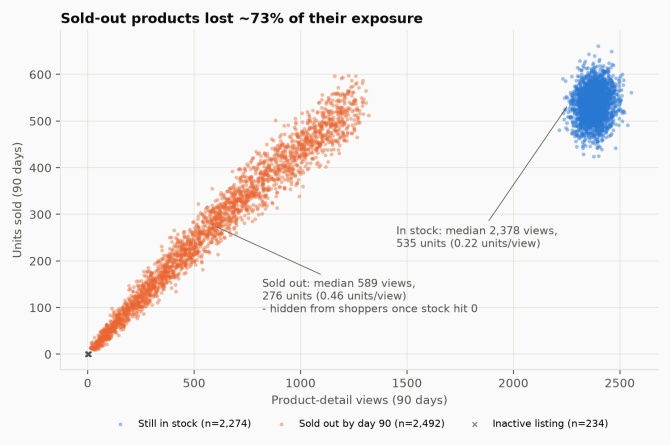


Fig. 3. Two populations: in-stock products (≈2,400 views) and sold-out products, whose exposure collapsed (median 589 views).

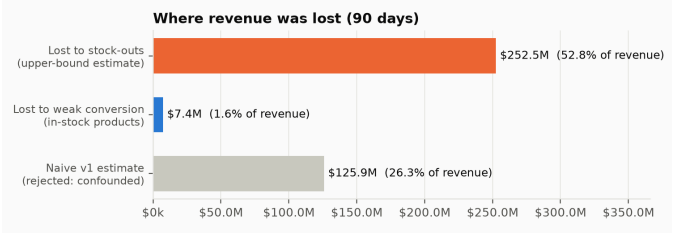


Fig. 4. Stock-outs cost ≈34× more than weak conversion. The rejected v1 figure is shown for contrast.

Result. Estimated revenue lost to stock-outs is **$252.5M**, an upper bound equal to 52.8% of realised revenue. It assumes a sold-out product would have attracted and converted traffic like its in-stock peers and could have been restocked immediately. Weak conversion among in-stock products costs $7.4M (1.6%). Every category has 41–58% of its products sold out, so the problem is systemic (initial stock is drawn uniformly from 10–1,000 units against ≈500 units of demand per product), not category-specific. Inactive listings received only 322 views in total and do not matter.

6. Technology selection, with measurements

The same two questions were run on all three engines (`integration/benchmark_engines.py`, best warm run; answers verified identical):

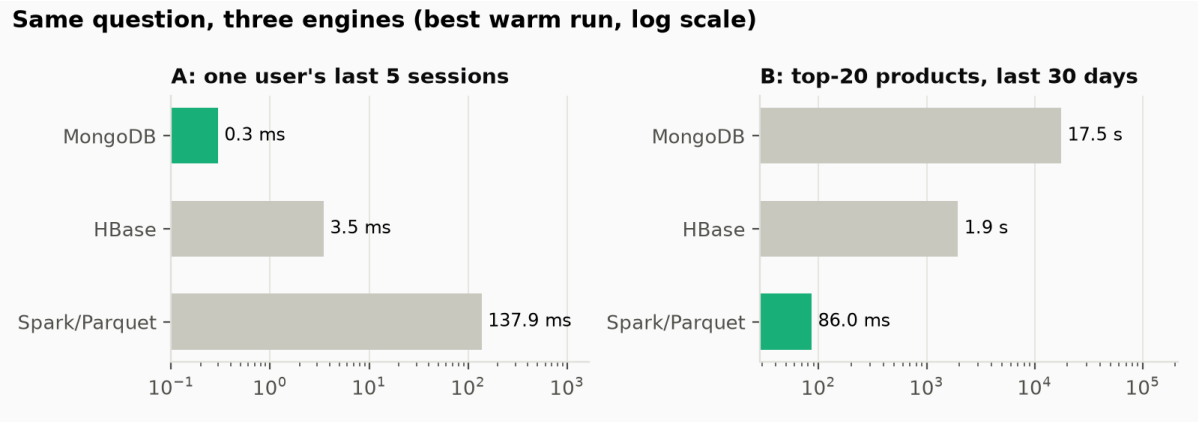


Fig. 5. Each engine wins its own access pattern by 1–2 orders of magnitude. HBase runs under amd64 emulation here, so its numbers are pessimistic.

For a point query (one user's latest sessions), MongoDB's B-tree index is fastest (0.3 ms) and HBase's row-key prefix scan is close behind (3.5 ms). Spark must scan a whole table (138 ms) and is the wrong tool. For an analytical question over 30 days of all products, Spark with Parquet partition pruning takes 86 ms, while MongoDB must unwind 664k documents (17.5 s). HBase must full-scan `product_daily` because its key is product-first (1.9 s). A date-first key would fix that but would create a write hotspot on “today”. The row key can serve one dominant access pattern well, and ours serves the per-product history. These measurements are the

justification for the architecture: **MongoDB for entities, HBase for keyed time series, Spark for scans and joins**.

7. Scalability

- **Spark**: every job is expressed as DataFrame/RDD transformations without driver-side loops, except two deliberately small driver reads (the 5k-entry broadcast and the shuffle-test permutation, which at 100× scale should become a distributed `rand()`-keyed shuffle). On a cluster, the JSON → Parquet step would read from HDFS/S3, with partitions sized at about 128 MB.
- **MongoDB**: shard `transactions` and `sessions` on **hashed** `user_id`. That gives even distribution, and per-user queries still target one shard. A **ranged** key on time would favour time-window analytics but hotspot inserts. The cross-journey aggregation (P3) does not scale on any shard key and belongs in Spark.
- **HBase**: regions split automatically. User ids generated sequentially would hotspot on load, which we mitigated by pre-splitting. At 50× scale a salt prefix (`hash(user)%N`) would be added. Enabling Snappy/ZSTD compression (not available in this image) would cut the 184 MB per 200k sessions by roughly 3–4×.
- **Integration**: replace driver-side reads with the MongoDB Spark Connector and HBase-Spark bulk get/scan, so that reads are parallel and locality-aware.

8. Key findings and business insights

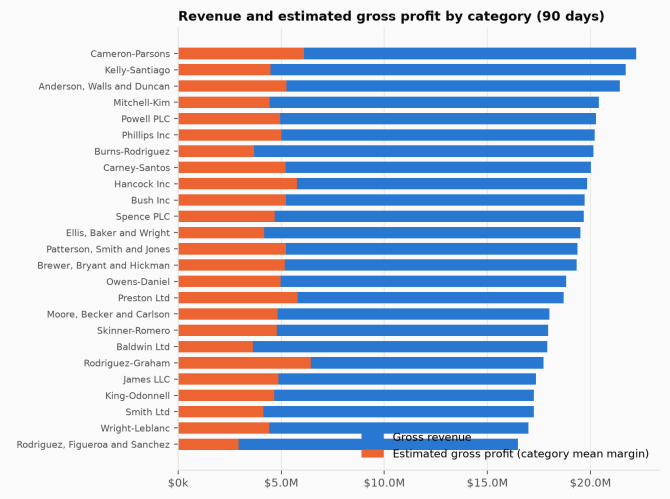


Fig. 6. Revenue is spread evenly ($16.5–22.2M per category), but estimated margins range from 17.8% to 36.3%, so profit ranking ≠ revenue ranking (Spearman 0.34). Kelly-Santiago is #2 by revenue but #18 by profit.

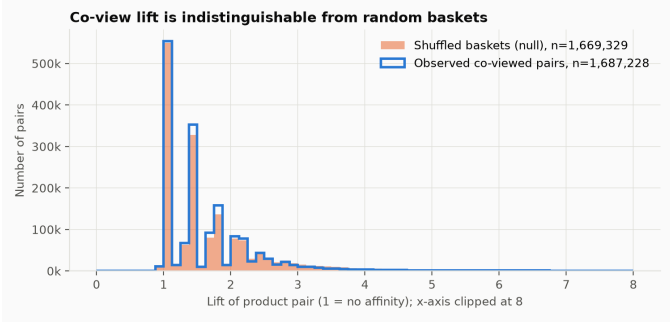


Fig. 7a. Co-view lift for 1.69M pairs matches randomly shuffled baskets (median 1.43 vs 1.43). The highest “lifts” (up to 560) are pairs seen 3 times.

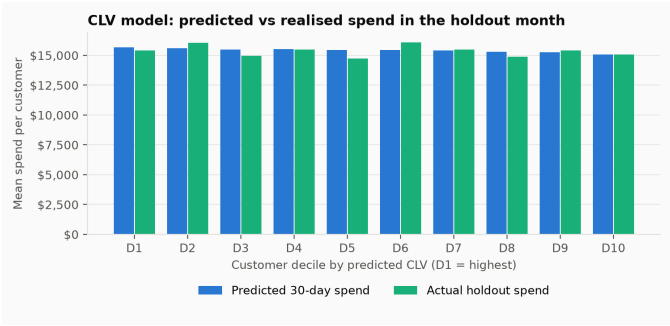


Fig. 7b. Predicted CLV deciles vs realised holdout spend: flat.

1. **Fix inventory before anything else.** Half the catalogue sold out, and sold-out products lost 73% of their exposure. Stock-outs are the single largest revenue lever (≈$250M upper bound vs $7M from conversion). Action: reorder points from daily sell-through (available per product in HBase `product_daily`).
2. **Measure the ~58% of revenue that has no session behind it.** Attribution and journey analytics can explain less than half of revenue until those orders are linked.
3. **Treat “converted” labels and historical prices as unreliable** (149k false conversions; 27.9% of lines at the wrong price) until the pipeline is fixed.
4. **Don't ship recommendations or CLV tiers built on this data.** The shuffle test and the holdout show there is no signal. With real data the same pipelines apply unchanged, and the validation steps will show

whether signal exists.

5. **Manage categories on profit, not revenue.** With margins of 17.8–36.3%, the revenue leaderboard misleads: Kelly-Santiago ranks #2 by revenue but #18 by estimated profit. (Caveat: margins are category means, see §2.)
6. **No channel, device or cohort effects.** Marketing spend cannot be allocated on these dimensions. Every cohort places 0.55 orders/member/day (±6%) once partial months are normalised by days observed.

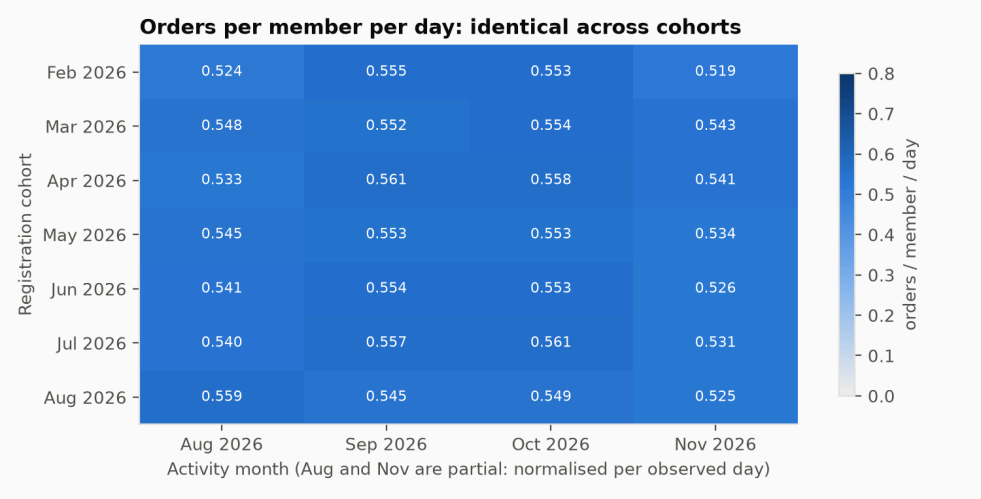


Fig. 8. Cohorts, normalised per observed day (raw monthly figures swung ±65% purely because August and November are partial months).

9. Limitations and future work

- **Synthetic data.** Users, products and timestamps are drawn independently. Several analyses therefore return well-supported *negative* results; they demonstrate method, not real customer behaviour.
- **Single-node deployment.** Spark local mode, one MongoDB instance and standalone HBase under x86 emulation. Absolute timings are indicative; the orders-of-magnitude differences are what matter.
- The stock-out loss is an **upper bound** under stated assumptions. Margins are category-level estimates because products lack subcategories. CLV ignores discounting and churn.
- HBase holds a 10% user sample of sessions. The CLV model therefore uses 1,000 users.
- **Future work:** Structured Streaming from a message broker (e.g. MQTT/Kafka) into HBase counters for real-time product metrics; MongoDB change streams to keep `purchase_summary` current; the Mongo/HBase Spark connectors; BG/NBD + Gamma-Gamma CLV on real data; collaborative filtering (ALS), but only after a lift-vs-null test shows there is signal to learn.

Reproducibility. `./run_all.sh` rebuilds everything (≈15 min). Session and transaction IDs use unseeded `uuid4`, so IDs differ between runs; all other values are seeded. Every figure and number in this report comes from the files in `output/` and `logs/` of the repository.

Use of AI assistance. An AI assistant (Claude, Anthropic) was used to help write and debug the code and to draft this report. All code was executed on the author's machine, and every result was produced by those runs and checked against the output files.